

Part III

High-Dimensional Systems

Chapter 17

Vector Spaces, Span, and Basis

17.1 Vector Spaces

Vector spaces are collections of vectors. The most common spaces are \mathbb{R}^2 , \mathbb{R}^3 , and \mathbb{R}^n — the spaces that include all 2-, 3-, and n -dimensional vectors. We can construct *subspaces* by specifying only a subset of the vectors in a space. For example, the set of all 3-dimensional vectors with only integer entries is a subspace of \mathbb{R}^3 .

Remember that \mathbb{R}^2 is not a subspace of \mathbb{R}^3 ; they are completely separate, non-overlapping spaces.

17.2 Span

A set of m vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m$ is said to *span* a space V if any vector \mathbf{u} in V can be written as a linear combination of the vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m$. This is equivalent to saying there exist scalars a_1, a_2, \dots, a_m such that

$$\mathbf{u} = a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_m\mathbf{v}_m.$$

Writing a vector \mathbf{u} as a linear combination of $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m$ is called *decomposing* \mathbf{u} over $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_m$. If a set of vectors spans a space, they can be used to decompose any other vector in the space.

We've already seen vector composition using a special set of vectors $\hat{\mathbf{e}}_j$, the unit vectors with only one nonzero entry at element j . For example, the vector

$$\begin{pmatrix} -2 \\ 4 \\ 5 \end{pmatrix} = -2 \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} + 4 \begin{pmatrix} 0 \\ 1 \\ 0 \end{pmatrix} + 5 \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix}.$$

Thus the vectors $\hat{\mathbf{e}}_1, \hat{\mathbf{e}}_2, \hat{\mathbf{e}}_3$ spans \mathbb{R}^3 . In general, the set of vectors $\hat{\mathbf{e}}_1, \hat{\mathbf{e}}_2, \dots, \hat{\mathbf{e}}_n$ spans the space \mathbb{R}^n . Are these the only sets of vectors that span these spaces?

No, there are infinitely many sets of vectors that span each space. Consider the vectors $\begin{pmatrix} 1 \\ 1 \end{pmatrix}$ and $\begin{pmatrix} -1 \\ 1 \end{pmatrix}$. We can show that these vectors span \mathbb{R}^2 by showing that any vector \mathbf{u} in \mathbb{R}^2 can be written as a linear combination

$$\mathbf{u} = \begin{pmatrix} u_1 \\ u_2 \end{pmatrix} = a_1 \begin{pmatrix} 1 \\ 1 \end{pmatrix} + a_2 \begin{pmatrix} -1 \\ 1 \end{pmatrix}.$$

Finding the coefficients a_1 and a_2 is akin to solving the system of linear equations

$$\begin{aligned} a_1 - a_2 &= u_1 \\ a_1 + a_2 &= u_2 \end{aligned}$$

which has the unique solution

$$a_1 = \frac{u_1 + u_2}{2}, \quad a_2 = \frac{u_2 - u_1}{2}.$$

To demonstrate, let $\mathbf{u} = \begin{pmatrix} -2 \\ 4 \end{pmatrix}$. Then $a_1 = 1$ and $a_2 = 3$, so

$$a_1 \begin{pmatrix} 1 \\ 1 \end{pmatrix} + a_2 \begin{pmatrix} -1 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 \\ 1 \end{pmatrix} + 3 \begin{pmatrix} -1 \\ 1 \end{pmatrix} = \begin{pmatrix} 1 - 3 \\ 1 + 3 \end{pmatrix} = \begin{pmatrix} -2 \\ 4 \end{pmatrix}.$$

We've shown that there are at least two sets of vectors that span \mathbb{R}^2 , $\left\{ \begin{pmatrix} 1 \\ 0 \end{pmatrix}, \begin{pmatrix} 0 \\ 1 \end{pmatrix} \right\}$ and $\left\{ \begin{pmatrix} 1 \\ 1 \end{pmatrix}, \begin{pmatrix} -1 \\ 1 \end{pmatrix} \right\}$. How can we say there are infinitely many? If vectors \mathbf{v}_1 and \mathbf{v}_2 span a space V , then the vectors $k_1\mathbf{v}_1$ and $k_2\mathbf{v}_2$ also span V for any scalars k_1 and k_2 . To prove this, remember that any vector \mathbf{u} can be decomposed onto \mathbf{v}_1 and \mathbf{v}_2 , i.e.

$$\begin{aligned} \mathbf{u} &= a_1\mathbf{v}_1 + a_2\mathbf{v}_2 \\ &= \frac{a_1}{k_1}(k_1\mathbf{v}_1) + \frac{a_2}{k_2}(k_2\mathbf{v}_2) \end{aligned}$$

Since (a_1/k_1) and (a_2/k_2) are simply scalars, we've shown that \mathbf{u} can be decomposed onto the vectors $k_1\mathbf{v}_1$ and $k_2\mathbf{v}_2$. Therefore, $k_1\mathbf{v}_1$ and $k_2\mathbf{v}_2$ must also span \mathbb{R}^2 . For example, the vectors $\begin{pmatrix} 3 \\ 0 \end{pmatrix}$ and $\begin{pmatrix} 0 \\ -1/2 \end{pmatrix}$ are scalar multiples of $\begin{pmatrix} 1 \\ 0 \end{pmatrix}$ and $\begin{pmatrix} 0 \\ 1 \end{pmatrix}$. The

Since k_1 and k_2 are arbitrary, this allows us to generate infinitely many sets of vectors that span any space from a single spanning set.

former two vectors must therefore span \mathbb{R}^2 , so we can decompose the vector $\begin{pmatrix} -2 \\ 4 \end{pmatrix}$ onto them:

$$\begin{pmatrix} -2 \\ 4 \end{pmatrix} = -\frac{2}{3} \begin{pmatrix} 3 \\ 0 \end{pmatrix} - 8 \begin{pmatrix} 0 \\ 1/2 \end{pmatrix}.$$

Similarly, if \mathbf{v}_1 and \mathbf{v}_2 span a space V , the vectors \mathbf{v}_1 and $(\mathbf{v}_1 + \mathbf{v}_2)$ also span V :

$$\begin{aligned} \mathbf{u} &= a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 \\ &= a_1 \mathbf{v}_1 + a_2 (\mathbf{v}_1 + \mathbf{v}_2) - a_2 \mathbf{v}_1 \\ &= (a_1 - a_2) \mathbf{v}_1 + a_2 (\mathbf{v}_1 + \mathbf{v}_2) \end{aligned}$$

If scalars a_1 and a_2 decompose \mathbf{u} over \mathbf{v}_1 and \mathbf{v}_2 , then $(a_1 - a_2)$ and a_2 decompose \mathbf{u} over \mathbf{v}_1 and $(\mathbf{v}_1 + \mathbf{v}_2)$.

17.3 Review: Linear Independence

We said before that vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ are *linearly independent* if and only if

$$a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_n \mathbf{v}_n = \mathbf{0}$$

implies that all coefficients a_1, a_2, \dots, a_n are zero. No linear combination of a set of linearly independent vectors can be the zero vector except for the trivial case where all the coefficients are zero. We often say that a set of vectors are linearly dependent if one of the vectors can be written as a linear combination of the others, i.e.

$$\mathbf{v}_i = a_1 \mathbf{v}_1 + \dots + a_{i-1} \mathbf{v}_{i-1} + a_{i+1} \mathbf{v}_{i+1} + \dots + a_n \mathbf{v}_n.$$

Moving the vector \mathbf{v}_i to the right hand side

$$\mathbf{0} = a_1 \mathbf{v}_1 + \dots + a_{i-1} \mathbf{v}_{i-1} - \mathbf{v}_i + a_{i+1} \mathbf{v}_{i+1} + \dots + a_n \mathbf{v}_n$$

we see 1.) a linear combination of the vectors sums to the zero vector on the left, and 2.) at least one of the coefficients (the -1 in front of \mathbf{v}_i) is nonzero. This is consistent with are above definition of linear independence. We said these vectors were not linearly independent, so it is possible for a linear combination to sum to zero using at least one nonzero coefficient.

17.4 Basis

The concepts of span and linear independence are a powerful combination. Any linearly independent set of vectors that span a space V are called a *basis* for V .

Any vector in a space be decomposed over a set of vectors that span the space. **However, every vector in a space has a unique decomposition over an associated basis.** Said another way, for every vector in a space, there are only one set of coefficients a_1, a_2, \dots, a_n that decompose it over a basis $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$.

We can prove that a decomposition over a basis is unique by contradiction. Suppose there were two sets of coefficients – a_1, a_2, \dots, a_n and b_1, b_2, \dots, b_n – that decomposed a vector \mathbf{u} over a basis $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$. Then

$$\mathbf{u} = a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_n\mathbf{v}_n = b_1\mathbf{v}_1 + b_2\mathbf{v}_2 + \dots + b_n\mathbf{v}_n.$$

We can move all the right hand side over to the left and group terms to give

$$(a_1 - b_1)\mathbf{v}_1 + (a_2 - b_2)\mathbf{v}_2 + \dots + (a_n - b_n)\mathbf{v}_n = \mathbf{0}.$$

Remember that $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ is a basis, so the vectors must be linearly independent. By the definition of linear independence, the only way the above equation can be true is if all the coefficients are zero. This implies that $a_1 = b_1$, $a_2 = b_2$, and so on. Clearly this is a violation of our original statement that a_1, a_2, \dots, a_n and b_1, b_2, \dots, b_n were different. Therefore, there can only be one way to decompose any vector onto a basis.

17.4.1 Testing if vectors form a basis

Every basis for a vector space has the same number of vectors. This number is called the *dimension* of the vector space. For standard vectors spaces like \mathbb{R}^n , the dimension is n . The dimension of \mathbb{R}^2 is 2, and the dimension of \mathbb{R}^3 is 3.

Any set of vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ is a basis for a space V if and only if:

1. The number of vectors (n) matches the dimension of V .
2. The vectors span V .
3. The vectors are linearly independent.

Proving any two of the above statements automatically implies the third is true.

We get to choose which two of the above three statements to prove when verifying that $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ is a basis. The first statement is usually trivial — does the number of vectors match the dimension? We almost always choose to prove the first statement. Proving that vectors are linearly independent is always easier than proving the vectors span the space. If we collect the vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ into a matrix, the rank of this matrix should be n if the vectors are linearly independent.

Most people think of dimension as the number of elements in a vector. While the true definition of dimension is the number of vectors in the basis, counting elements in a vector works for spaces like \mathbb{R}^n . To see why, remember that the Cartesian unit vectors $\hat{\mathbf{e}}_i$ form a basis for \mathbb{R}^n , but we need n of these vectors, one per element.

17.4.2 Decomposing onto a basis

How do we decompose a vector onto a basis? Remember that decomposing \mathbf{u} over $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ is equivalent to finding a set of coefficients a_1, a_2, \dots, a_n such that

$$a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_n\mathbf{v}_n = \mathbf{u}.$$

Let's collect the vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ into a matrix \mathbf{V} , where each vector \mathbf{v}_i is the i th column in \mathbf{V} . Then

$$(\mathbf{v}_1 \quad \mathbf{v}_2 \quad \dots \quad \mathbf{v}_n) \begin{pmatrix} a_1 \\ a_2 \\ \vdots \\ a_n \end{pmatrix} = \mathbf{V}\mathbf{a} = \mathbf{u}.$$

We see that finding the coefficients a_1, a_2, \dots, a_n that decompose a vector \mathbf{u} onto a basis $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ is equivalent to solving the linear system $\mathbf{V}\mathbf{a} = \mathbf{u}$.

By formulating vector decomposition as a linear system, we can easily see why the decomposition of a vector over a basis is unique. In \mathbb{R}^n , the basis contains n vectors, each with n elements. So, the matrix \mathbf{V} is a square, $n \times n$ matrix. Since the vectors in the basis (and therefore the columns in \mathbf{V}) are linearly independent, the matrix \mathbf{V} has full rank. Thus, the solution to $\mathbf{V}\mathbf{a} = \mathbf{u}$ must be unique, implying that the decomposition of every vector onto a basis is unique.

Since \mathbf{V} is square and full rank, its inverse (\mathbf{V}^{-1}) must exist. The system has a unique solution $\mathbf{a} = \mathbf{V}^{-1}\mathbf{u}$.

17.5 Orthogonal and Orthonormal Vectors

A set of vectors is an *orthogonal set* if every vector in the set is orthogonal to every other vector in the set. If every vector in an orthogonal set has been normalized, we say the vectors form an *orthonormal set*. Orthogonal and orthonormal sets are ideal candidates for basis vectors. Since there is no “overlap” among the vectors, we can easily decompose other vectors onto orthogonal basis vectors.

Imagine you have an orthogonal set of vectors you want to use as a basis. We'll assume you have the correct number of vectors (equal to the dimension of your space) for this to be possible. Based on the above requirements for basis vectors, we need only to show that these vectors are linearly independent. If so, they are a basis. For a set of n orthogonal vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$, linear independence requires that

$$a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_n\mathbf{v}_n = \mathbf{0}$$

if and only if all the coefficients a_1, a_2, \dots, a_n are equal to zero. Let's take the dot product of both sides of the above equation with the vector \mathbf{v}_1 :

$$(a_1\mathbf{v}_1 + a_2\mathbf{v}_2 + \dots + a_n\mathbf{v}_n) \cdot \mathbf{v}_1 = \mathbf{0} \cdot \mathbf{v}_1.$$

On the right hand side, we know that $\mathbf{0} \cdot \mathbf{v}_1 = 0$ for any vector \mathbf{v}_1 . We also distribute the dot product on the left hand side to give

$$a_1 \mathbf{v}_1 \cdot \mathbf{v}_1 + a_2 \mathbf{v}_2 \cdot \mathbf{v}_1 + \cdots + a_n \mathbf{v}_n \cdot \mathbf{v}_1 = 0.$$

Since all the vectors are orthogonal, $\mathbf{v}_i \cdot \mathbf{v}_1$ is zero except when $i = 1$. Canceling out all the dot products equal to zero shows that

$$a_1 \mathbf{v}_1 \cdot \mathbf{v}_1 = a_1 \|\mathbf{v}_1\|^2 = 0.$$

We know that $\|\mathbf{v}_1\|^2 \neq 0$, so the only way $a_1 \mathbf{v}_1 \cdot \mathbf{v}_1$ can equal zero is if a_1 is zero. If we repeat this entire process by taking the dot product with \mathbf{v}_2 instead of \mathbf{v}_1 , we will find that $a_2 = 0$. This continues with $\mathbf{v}_3, \dots, \mathbf{v}_n$ until we can say that $a_1 = a_2 = \cdots = a_n = 0$. Therefore, if the vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ are an orthogonal (or orthonormal) set, they are linearly independent.

17.5.1 Decomposing onto orthonormal vectors

We saw previously that finding the coefficients to decompose a vector onto a basis requires solving a system of linear equations. For high-dimensional spaces, solving such a system can be computationally expensive. Fortunately, decomposing a vector onto an orthonormal basis is far more efficient.

Theorem. *The decomposition of a vector \mathbf{u} onto an orthonormal basis $\hat{\mathbf{v}}_1, \hat{\mathbf{v}}_2, \dots, \hat{\mathbf{v}}_n$ given by*

$$\mathbf{u} = a_1 \hat{\mathbf{v}}_1 + a_2 \hat{\mathbf{v}}_2 + \cdots + a_n \hat{\mathbf{v}}_n$$

has coefficients

$$a_1 = \mathbf{u} \cdot \hat{\mathbf{v}}_1$$

$$a_2 = \mathbf{u} \cdot \hat{\mathbf{v}}_2$$

$$\vdots$$

$$a_n = \mathbf{u} \cdot \hat{\mathbf{v}}_n$$

Proof. We use a similar strategy as when we proved the linear independence of orthogonal sets. Let's start with the formula for vector decomposition

$$\mathbf{u} = a_1 \hat{\mathbf{v}}_1 + a_2 \hat{\mathbf{v}}_2 + \cdots + a_n \hat{\mathbf{v}}_n.$$

Taking the dot product of both sides with the vector $\hat{\mathbf{v}}_1$ yields (after distributing the right hand side)

$$\mathbf{u} \cdot \hat{\mathbf{v}}_1 = a_1 \hat{\mathbf{v}}_1 \cdot \hat{\mathbf{v}}_1 + a_2 \hat{\mathbf{v}}_2 \cdot \hat{\mathbf{v}}_1 + \cdots + a_n \hat{\mathbf{v}}_n \cdot \hat{\mathbf{v}}_1.$$

We use the symbol $\hat{\mathbf{v}}_i$ for vectors in an orthonormal set to remind us that each vector has been normalized.

Because all the vectors $\hat{\mathbf{v}}_1, \hat{\mathbf{v}}_2, \dots, \hat{\mathbf{v}}_n$ are orthogonal, the only nonzero term on the right hand side is $a_1 \hat{\mathbf{v}}_1 \cdot \hat{\mathbf{v}}_1$, so

$$\mathbf{u} \cdot \hat{\mathbf{v}}_1 = a_1 \hat{\mathbf{v}}_1 \cdot \hat{\mathbf{v}}_1.$$

By definition of the dot product, $\hat{\mathbf{v}}_1 \cdot \hat{\mathbf{v}}_1 = \|\hat{\mathbf{v}}_1\|^2$. Since $\hat{\mathbf{v}}_1$ is a unit vector, $\|\hat{\mathbf{v}}_1\|^2 = 1$. Thus $a_1 = \mathbf{u} \cdot \hat{\mathbf{v}}_1$. By repeating the same procedure with $\hat{\mathbf{v}}_2$, we find that $a_2 = \mathbf{u} \cdot \hat{\mathbf{v}}_2$, and so on. \square

Decomposing vectors over an orthonormal basis is efficient, requiring only a series of dot products to compute the coefficients. For example, we can decompose the

vector $\mathbf{u} = \begin{pmatrix} 7 \\ -5 \\ 10 \end{pmatrix}$ over the orthonormal basis $\left\{ \hat{\mathbf{v}}_1 = \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix}, \hat{\mathbf{v}}_2 = \begin{pmatrix} 0 \\ 3/5 \\ 4/5 \end{pmatrix}, \hat{\mathbf{v}}_3 = \begin{pmatrix} 0 \\ 4/5 \\ -3/5 \end{pmatrix} \right\}$.

$$a_1 = \mathbf{u} \cdot \hat{\mathbf{v}}_1 = \begin{pmatrix} 7 \\ -5 \\ 10 \end{pmatrix} \cdot \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} = 7 + 0 + 0 = 7$$

$$a_2 = \mathbf{u} \cdot \hat{\mathbf{v}}_2 = \begin{pmatrix} 7 \\ -5 \\ 10 \end{pmatrix} \cdot \begin{pmatrix} 0 \\ 3/5 \\ 4/5 \end{pmatrix} = 0 - 3 + 8 = 5$$

$$a_3 = \mathbf{u} \cdot \hat{\mathbf{v}}_3 = \begin{pmatrix} 7 \\ -5 \\ 10 \end{pmatrix} \cdot \begin{pmatrix} 0 \\ 4/5 \\ -3/5 \end{pmatrix} = 0 - 4 - 6 = -10$$

The decomposition of \mathbf{u} is

$$\mathbf{u} = 7 \begin{pmatrix} 1 \\ 0 \\ 0 \end{pmatrix} + 5 \begin{pmatrix} 0 \\ 3/5 \\ 4/5 \end{pmatrix} - 10 \begin{pmatrix} 0 \\ 4/5 \\ -3/5 \end{pmatrix} = \begin{pmatrix} 7 + 0 + 0 \\ 0 + 3 - 8 \\ 0 + 4 + 6 \end{pmatrix} = \begin{pmatrix} 7 \\ -5 \\ 10 \end{pmatrix}.$$

17.5.2 Checking an orthonormal set

Given a set of vectors $\hat{\mathbf{v}}_1, \hat{\mathbf{v}}_2, \dots, \hat{\mathbf{v}}_n$, how can we verify that they are orthonormal? We need to test two things.

1. All vectors are normalized ($\|\hat{\mathbf{v}}_i\| = 1$ for all $\hat{\mathbf{v}}_i$).
2. All vectors are mutually orthogonal ($\hat{\mathbf{v}}_i \cdot \hat{\mathbf{v}}_j = 0$ for all $i \neq j$).

The first test is straightforward. The second can be a little cumbersome, as we need to test all $n^2 - n/2$ pairs of vectors for orthogonality. A simpler, albeit more sometimes more computationally intensive approach, is to collect the vectors into a matrix $\mathbf{V} = (\mathbf{v}_1 \ \mathbf{v}_2 \ \cdots \ \mathbf{v}_n)$. Then the set of vectors is orthonormal if and only if $\mathbf{V}^{-1} = \mathbf{V}^T$. While inverting a matrix is “expensive,” for small to medium size vector sets this method avoids the need to iterate over all pairs of vectors to test for orthonormality.

17.5.3 Projections

Our next goal will be to create orthonormal sets of vectors from sets that are not orthogonal. Before introducing such an algorithm, we need to develop a geometric tool — the vector *projection*. The projection of vector \mathbf{v} onto vector \mathbf{u} is a vector that points along \mathbf{u} with length equal to the “shadow” of \mathbf{v} onto \mathbf{u} . Previously we used the dot product to calculate the magnitude of the projection of \mathbf{v} onto \mathbf{u} , which was a scalar equal to $\|\mathbf{v}\| \cos \theta$, where θ is the angle between \mathbf{v} and \mathbf{u} . To calculate the actual projection, we multiply the magnitude of the projection ($\|\mathbf{v}\| \cos \theta$) by a unit vector that points along \mathbf{u} . Thus the projection of \mathbf{v} onto \mathbf{u} is defined as

$$\text{proj}_{\mathbf{u}}(\mathbf{v}) = (\|\mathbf{v}\| \cos \theta) \hat{\mathbf{u}}.$$

By definition, $\hat{\mathbf{u}} = \mathbf{u}/\|\mathbf{u}\|$. Also, we note that $\mathbf{v} \cdot \mathbf{u} = \|\mathbf{v}\| \|\mathbf{u}\| \cos \theta$, so the expression $\|\mathbf{u}\| \cos \theta$ can be written in terms of the dot product $(\mathbf{v} \cdot \mathbf{u})/\|\mathbf{u}\|$. We can rewrite our formula for the projection using only dot products:

$$\text{proj}_{\mathbf{u}}(\mathbf{v}) = (\|\mathbf{v}\| \cos \theta) \hat{\mathbf{u}} = \frac{\mathbf{v} \cdot \mathbf{u}}{\|\mathbf{u}\|} \frac{\mathbf{u}}{\|\mathbf{u}\|} = \frac{\mathbf{v} \cdot \mathbf{u}}{\|\mathbf{u}\|^2} \mathbf{u} = \frac{\mathbf{v} \cdot \mathbf{u}}{\mathbf{u} \cdot \mathbf{u}} \mathbf{u}.$$

We can use the projection to make any two vectors orthogonal, as demonstrated by the following theorem.

Theorem. *Given any vectors \mathbf{v} and \mathbf{u} , the vector $\mathbf{v} - \text{proj}_{\mathbf{u}}(\mathbf{v})$ is orthogonal to \mathbf{u} .*

Proof. If the vector $\mathbf{v} - \text{proj}_{\mathbf{u}}(\mathbf{v})$ is orthogonal to \mathbf{u} , the dot product between these vectors must be zero.

$$\begin{aligned} (\mathbf{v} - \text{proj}_{\mathbf{u}}(\mathbf{v})) \cdot \mathbf{u} &= \left(\mathbf{v} - \frac{\mathbf{v} \cdot \mathbf{u}}{\mathbf{u} \cdot \mathbf{u}} \mathbf{u} \right) \cdot \mathbf{u} \\ &= \mathbf{v} \cdot \mathbf{u} - \frac{\mathbf{v} \cdot \mathbf{u}}{\mathbf{u} \cdot \mathbf{u}} \mathbf{u} \cdot \mathbf{u} \\ &= \mathbf{v} \cdot \mathbf{u} - \mathbf{v} \cdot \mathbf{u} \\ &= 0 \end{aligned}$$

Interestingly, the proof of this method (not shown here) reveals that if the columns of \mathbf{V} are an orthonormal set, the rows of \mathbf{V} are also an orthonormal set!




Figure 17.1: The projection is a vector “shadow” of one vector onto another.

□

Subtracting the projection of \mathbf{v} onto \mathbf{u} from the vector \mathbf{v} “corrects” \mathbf{v} by removing its overlap with \mathbf{u} . The resulting vector is a vector closest to \mathbf{v} that is still orthogonal to \mathbf{u} .

17.5.4 Creating orthonormal basis vectors

We can make any two vectors orthogonal by adjusting one based on its projection onto the other. We can apply these corrections iteratively to make any set of linearly independent vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$ into an orthonormal basis set $\hat{\mathbf{u}}_1, \hat{\mathbf{u}}_2, \dots, \hat{\mathbf{u}}_n$. First, we set

$$\mathbf{u}_1 = \mathbf{v}_1.$$

We leave this first vector unchanged. All other vectors will be made orthogonal to it (and each other). Next, we create \mathbf{u}_2 by making \mathbf{v}_2 orthogonal to \mathbf{u}_1 :

$$\mathbf{u}_2 = \mathbf{v}_2 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_2).$$

Now we have two orthogonal vectors, \mathbf{u}_1 and \mathbf{u}_2 . We continue by creating \mathbf{u}_3 from \mathbf{v}_3 , but this time we must make \mathbf{v}_3 orthogonal to both \mathbf{u}_1 and \mathbf{u}_2 :

$$\mathbf{u}_3 = \mathbf{v}_3 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_3) - \text{proj}_{\mathbf{u}_2}(\mathbf{v}_3).$$

We continue this process for all n vectors, making each vector \mathbf{v}_i orthogonal to all the newly created orthogonal vectors $\mathbf{u}_1, \dots, \mathbf{u}_{i-1}$. This approach is called the Gram-Schmidt algorithm. Given a set of vectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$, we create a set of orthogonal vectors

$$\begin{aligned} \mathbf{u}_1 &= \mathbf{v}_1 \\ \mathbf{u}_2 &= \mathbf{v}_2 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_2) \\ \mathbf{u}_3 &= \mathbf{v}_3 - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_3) - \text{proj}_{\mathbf{u}_2}(\mathbf{v}_3) \\ &\vdots \\ \mathbf{u}_i &= \mathbf{v}_i - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_i) - \dots - \text{proj}_{\mathbf{u}_{i-1}}(\mathbf{v}_i) \\ &\vdots \\ \mathbf{u}_n &= \mathbf{v}_n - \text{proj}_{\mathbf{u}_1}(\mathbf{v}_n) - \dots - \text{proj}_{\mathbf{u}_{n-1}}(\mathbf{v}_n) \end{aligned}$$

The Gram-Schmidt algorithm produces an orthogonal set of vectors. To make the set orthonormal, we must subsequently normalize each vector.




Figure 17.2: Subtracting the projection of \mathbf{v} onto \mathbf{u} from \mathbf{v} makes the vectors orthogonal.

We must begin with linearly independent vectors. Otherwise, orthogonalization will turn one of the vectors into the zero vector, which is not allowed in a basis.

Said more succinctly,

$$\mathbf{u}_i = \mathbf{v}_i - \sum_{k=1}^{i-1} \text{proj}_{\mathbf{u}_k}(\mathbf{v}_i)$$

Chapter 18

Eigenvectors and Eigenvalues

Consider the matrix

$$\mathbf{A} = \begin{pmatrix} 2 & 7 \\ -1 & -6 \end{pmatrix}.$$

Multiplying \mathbf{A} by the vector $\mathbf{x}_1 = \begin{pmatrix} -1 \\ 1 \end{pmatrix}$ gives an interesting result.

$$\mathbf{A}\mathbf{x}_1 = \begin{pmatrix} 2 & 7 \\ -1 & -6 \end{pmatrix} \begin{pmatrix} -1 \\ 1 \end{pmatrix} = \begin{pmatrix} 5 \\ -5 \end{pmatrix} = -5 \begin{pmatrix} -1 \\ 1 \end{pmatrix} = -5\mathbf{x}_1.$$

Similarly, with $\mathbf{x}_2 = \begin{pmatrix} -7 \\ 1 \end{pmatrix}$:

$$\mathbf{A}\mathbf{x}_2 = \begin{pmatrix} 2 & 7 \\ -1 & -6 \end{pmatrix} \begin{pmatrix} -7 \\ 1 \end{pmatrix} = \begin{pmatrix} -7 \\ 1 \end{pmatrix} = \mathbf{x}_2.$$

In both cases, multiplication by \mathbf{A} returned a scalar multiple of the vector (-5 for \mathbf{x}_1 and 1 for \mathbf{x}_2). This is not a property of solely the matrix \mathbf{A} , since the vector $\mathbf{x}_3 = \begin{pmatrix} 1 \\ 1 \end{pmatrix}$ is not transformed by a single scalar.

$$\mathbf{A}\mathbf{x}_3 = \begin{pmatrix} 2 & 7 \\ -1 & -6 \end{pmatrix} \begin{pmatrix} 1 \\ 1 \end{pmatrix} = \begin{pmatrix} 9 \\ 5 \end{pmatrix} \neq \lambda \mathbf{x}_3$$

Similarly, the results we are seeing are not properties of the vectors \mathbf{x}_1 and \mathbf{x}_2 , since they do not become scalar multiples of themselves when multiplied by other

matrices.

$$\mathbf{B} = \begin{pmatrix} 2 & 1 \\ -3 & 0 \end{pmatrix}$$

$$\mathbf{B}\mathbf{x}_1 = \begin{pmatrix} 2 & 1 \\ -3 & 0 \end{pmatrix} \begin{pmatrix} -1 \\ 1 \end{pmatrix} = \begin{pmatrix} -1 \\ 3 \end{pmatrix} \neq \lambda \mathbf{x}_1$$

$$\mathbf{B}\mathbf{x}_2 = \begin{pmatrix} 2 & 1 \\ -3 & 0 \end{pmatrix} \begin{pmatrix} -7 \\ 1 \end{pmatrix} = \begin{pmatrix} -13 \\ 21 \end{pmatrix} \neq \lambda \mathbf{x}_2$$

The phenomena we're observing is a result of the paring between the matrix \mathbf{A} and the vectors \mathbf{x}_1 and \mathbf{x}_2 . In general, we see that multiplying a vector by a matrix returns a scalar multiple of the vector, or

$$\mathbf{A}\mathbf{x} = \lambda \mathbf{x}.$$

Any vector \mathbf{x} that obeys the above relationship is called an *eigenvector* of the matrix \mathbf{A} . The scalar λ is called the *eigenvalue* associated with the eigenvector \mathbf{x} . The vector \mathbf{x} is an eigenvector of the matrix \mathbf{A} ; it is not generally an eigenvector of other matrices.

In the example above, the matrix $\mathbf{A} = \begin{pmatrix} 2 & 7 \\ -1 & -6 \end{pmatrix}$ has two eigenvectors, $\mathbf{v}_1 = \begin{pmatrix} -1 \\ 1 \end{pmatrix}$ with eigenvalue $\lambda_1 = -5$, and $\mathbf{v}_2 = \begin{pmatrix} -7 \\ 1 \end{pmatrix}$ with eigenvalue $\lambda_2 = 1$.

Eigenvectors were originally called *characteristic* vectors, as they describe the character of the matrix. German mathematicians dropped this nomenclature in favor of the German prefix "eigen-", which mean "own". An eigenvector can be viewed as one of a matrix's "own" vectors since it is not rotated when transformed by multiplication.

18.1 Properties of Eigenvectors and Eigenvalues

Only square matrices have eigenvectors and eigenvalues. To understand why the matrix must be square, remember that a non-square matrix with m rows and n columns transforms an n -dimensional vectors into an m -dimensional vector. Clearly, the m -dimensional output cannot be the n -dimensional input multiplied by a scalar!

An n by n matrix of real numbers can have up to n distinct eigenvectors. Each eigenvector is associated with an eigenvalue, although the eigenvalues can be duplicated. Said another way, two eigenvectors \mathbf{v}_1 and \mathbf{v}_2 of a matrix will never be the same, but the corresponding eigenvalues λ_1 and λ_2 can be identical.

Although the number of eigenvectors may vary, all eigenvectors for a matrix are linearly independent. Thus, if an n by n matrix has n eigenvectors, these vectors can be used as a basis (called an *eigenbasis*). If an eigenbasis exists for a matrix, decomposing vectors over this basis simplifies the process of matrix

An n by n matrix with n eigenvectors and n distinct eigenvalues is called a *perfect matrix*. As the name implies, perfect matrices are great to find, but somewhat uncommon.

multiplication. To illustrate, imagine we decompose the vector \mathbf{x} over a set of eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_n$. Decomposing \mathbf{x} means we can find coefficients a_1, \dots, a_n such that

$$\mathbf{x} = a_1\mathbf{v}_1 + \dots + a_n\mathbf{v}_n.$$

Now let's compute the product \mathbf{Ax} . We multiply both sides of the decomposition by \mathbf{A} .

$$\mathbf{Ax} = \mathbf{A}(a_1\mathbf{v}_1 + \dots + a_n\mathbf{v}_n)$$

We distribute the matrix \mathbf{A} into the sum on the right hand side and note that the constants a_i can be moved in front of the matrix multiplication.

$$\mathbf{Ax} = a_1\mathbf{Av}_1 + \dots + a_n\mathbf{Av}_n$$

Remember that $\mathbf{v}_1, \dots, \mathbf{v}_n$ are eigenvectors of \mathbf{A} , so $\mathbf{Av}_i = \lambda_i\mathbf{v}_i$. We can simplify the previous expression to

$$\mathbf{Ax} = a_1\lambda_1\mathbf{v}_1 + \dots + a_n\lambda_n\mathbf{v}_n.$$

We don't need to perform the multiplication at all! Instead, we can scale each eigenvector by the eigenvalue. Multiplying again by the matrix \mathbf{A} multiplies each eigenvector by its eigenvalue.

$$\mathbf{A}^2\mathbf{x} = a_1\lambda_1^2\mathbf{v}_1 + \dots + a_n\lambda_n^2\mathbf{v}_n$$

$$\mathbf{A}^k\mathbf{x} = a_1\lambda_1^k\mathbf{v}_1 + \dots + a_n\lambda_n^k\mathbf{v}_n$$

We use the notation \mathbf{A}^2 to denote \mathbf{AA} , \mathbf{A}^3 for \mathbf{AAA} , and \mathbf{A}^k for the product of k matrices \mathbf{A} .

18.2 Computing Eigenvectors and Eigenvalues

We can use the relationship between matrix multiplication and eigenvalues to find eigenvectors for any matrix. Our computational approach is based on the following theorem.

Theorem. *Given any (random) vector \mathbf{b} , repeated multiplication by the matrix \mathbf{A} will converge to the eigenvector of \mathbf{A} with the largest magnitude eigenvalue – provided the largest eigenvalue is unique. Said another way,*

$$\lim_{k \rightarrow \infty} \mathbf{A}^k \mathbf{b} = \mathbf{v}_{\max}.$$

Proof. We know that the product \mathbf{Ax} can be expressed as a linear combination of the eigenvectors and eigenvalues of \mathbf{A} , i.e. $\mathbf{Ax} = a_1\lambda_1\mathbf{v}_1 + \dots + a_n\lambda_n\mathbf{v}_n$. Thus

$$\lim_{k \rightarrow \infty} \mathbf{A}^k \mathbf{b} = \lim_{k \rightarrow \infty} (a_1\lambda_1^k\mathbf{v}_1 + \dots + a_n\lambda_n^k\mathbf{v}_n).$$

As k increases, the values λ_i^k grow very large. However, the λ_i do not grow at the same rate. The largest eigenvalue will grow the fastest. At very large values of k , the term associated with the largest eigenvalue will dominate the entire sum, so the result will point in only the direction of the associated eigenvector. Note that convergence to a single eigenvector requires that the largest eigenvalue be distinct. If two eigenvectors have the same (largest) eigenvalue, both terms in the above sum would “blow up” at the same rate. Repeated multiplications by \mathbf{A} would then converge to the sum of the two associated eigenvectors. \square

The above theorem allows us to find the eigenvector paired with the largest eigenvalue. While the direction of the eigenvector doesn’t change, its magnitude grows as the number of multiplication of \mathbf{A} increases. If convergence is slow, we might need to work with numbers before finding the eigenvector. To avoid numerical difficulties, we renormalize the vector after every multiplication by \mathbf{A} . This algorithm is called the Power Iteration method, which proceeds as follows:

1. Choose a random vector \mathbf{b}_0 . For fastest convergence, it helps to choose a vector close to \mathbf{v}_{\max} if possible. Normalize this vector to product $\hat{\mathbf{b}}_0 = \mathbf{b}_0 / \|\mathbf{b}_0\|$.
2. Compute vector $\mathbf{b}_1 = \mathbf{A}\hat{\mathbf{b}}_0$. Normalize this vector to give $\hat{\mathbf{b}}_1 = \mathbf{b}_1 / \|\mathbf{b}_1\|$.
3. Repeat step 2 to product $\hat{\mathbf{b}}_2, \hat{\mathbf{b}}_3, \dots, \hat{\mathbf{b}}_k$. Stop when all entries of $\hat{\mathbf{b}}_k$ do not change from the entries in $\hat{\mathbf{b}}_{k-1}$. The vector $\hat{\mathbf{b}}_k$ is an eigenvector of \mathbf{A} .

Now that we have the eigenvector \mathbf{v}_{\max} , how do we find the associated eigenvalue λ_{\max} ? We know that \mathbf{v}_{\max} is an eigenvector of \mathbf{A} , to $\mathbf{A}\mathbf{v}_{\max} = \lambda_{\max}\mathbf{v}_{\max}$. The i th element in $\mathbf{A}\mathbf{v}_{\max}$ should be equal to λ_{\max} times the i th element in \mathbf{v}_{\max} . However, since we only found a numerical approximation to the \mathbf{v}_{\max} , the estimate for λ_{\max} from each element in \mathbf{v}_{\max} might differ slightly. To “smooth out” these variations, compute the eigenvalue using the Rayleigh quotient:

$$\lambda_{\max} = \frac{\mathbf{v}_{\max} \cdot \mathbf{A}\mathbf{v}_{\max}}{\mathbf{v}_{\max} \cdot \mathbf{v}_{\max}}.$$

The dot product in the Rayleigh quotient averages out all of the small discrepancies between our estimate \mathbf{v}_{\max} and the true largest eigenvector. The Rayleigh quotient provides a numerically stable estimate of the largest eigenvalue.

Now that we’ve found the first eigenvector, how do we find the others? If we start the Power Iteration method over again using the matrix $(\mathbf{A} - \lambda_{\max}\mathbf{I})$ instead of \mathbf{A} , the algorithm will converge to the eigenvector associated with the second largest eigenvalue. We can subtract this eigenvalue from \mathbf{A} and repeat to find the

The eigenvector associated with the largest magnitude eigenvalue is called the *leading eigenvector*.

To see why the Rayleigh quotient works, consider an eigenvector \mathbf{v} for matrix \mathbf{A} with associated eigenvalue λ . Then

$$\frac{\mathbf{v} \cdot \mathbf{A}\mathbf{v}}{\mathbf{v} \cdot \mathbf{v}} = \frac{\mathbf{v} \cdot (\lambda\mathbf{v})}{\mathbf{v} \cdot \mathbf{v}} = \lambda \frac{\mathbf{v} \cdot \mathbf{v}}{\mathbf{v} \cdot \mathbf{v}} = \lambda.$$

third eigenvector, and so on. Proving Power Iteration is able to find subsequent eigenvectors is beyond the scope of this course. However, as we'll see later, finding only the first eigenvector is sufficient for addressing a number of interesting problems.

18.2.1 Eigenvalues and Eigenvectors in MATLAB

The MATLAB function `eig` computes eigenvalues and eigenvectors. The statement `[V,L] = eig(A)` involving an n by n matrix A returns two n by n matrices:

- Each column of the matrix V is an eigenvector A .
- The matrix L is a diagonal matrix. The i th entry on the diagonal is the eigenvalue associated with the i th column in V .

Remember that any vector that points in the same direction as an eigenvector of a matrix is also an eigenvector of that matrix. If the eigenvectors returned by computational systems like MATLAB are not what you expect, remember that they may be normalized or scaled – but still point along the same direction.

18.3 Applications

Eigenvalue and eigenvectors can be used to solve a number of interesting engineering and data science problems.

18.3.1 Solving Systems of ODEs

Consider the linear system of ODEs

$$\begin{aligned}\frac{dx_1}{dt} &= x_1 + 2x_2 \\ \frac{dx_2}{dt} &= 3x_1 + 2x_2\end{aligned}$$

with initial conditions $x_1(0) = 0$ and $x_2(0) = -4$. We can write this system using vectors and matrices as

$$\frac{d\mathbf{x}}{dt} = \mathbf{A}\mathbf{x}, \quad \mathbf{x}(0) = \mathbf{x}_0$$

where for the example above

$$\mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}, \quad \mathbf{A} = \begin{pmatrix} 1 & 2 \\ 3 & 2 \end{pmatrix}, \quad \mathbf{x}_0 = \begin{pmatrix} 0 \\ -4 \end{pmatrix}.$$

If we know the eigenvectors $\mathbf{v}_1, \dots, \mathbf{v}_n$ and eigenvalues $\lambda_1, \dots, \lambda_n$ for the matrix \mathbf{A} , we can compute the solution as

$$\mathbf{x}(t) = c_1 \mathbf{v}_1 e^{\lambda_1 t} + c_2 \mathbf{v}_2 e^{\lambda_2 t} + \dots + c_n \mathbf{v}_n e^{\lambda_n t}.$$

The scalars c_1, \dots, c_n are the constants of integration. To find these values, notice what happens to our solution at time $t = 0$:

$$\mathbf{x}(0) = \mathbf{x}_0 = c_1 \mathbf{v}_1 + c_2 \mathbf{v}_2 + \dots + c_n \mathbf{v}_n.$$

At $t = 0$, the right hand side is a decomposition of the initial conditions \mathbf{x}_0 . If we collect the eigenvectors as columns of a matrix $\mathbf{V} = (\mathbf{v}_1 \mathbf{v}_2 \dots \mathbf{v}_n)$, we can find the constants c_1, \dots, c_n by solving the linear system

$$\mathbf{V} \begin{pmatrix} c_1 \\ \vdots \\ c_n \end{pmatrix} = \mathbf{x}_0.$$

Returning to our original example, the matrix

$$\mathbf{A} = \begin{pmatrix} 1 & 2 \\ 3 & 2 \end{pmatrix}$$

has eigenvalue/eigenvector pairs

$$\lambda_1 = -1, \quad \mathbf{v}_1 = \begin{pmatrix} -1 \\ 1 \end{pmatrix} \quad \text{and} \quad \lambda_2 = 4, \quad \mathbf{v}_2 = \begin{pmatrix} 2 \\ 3 \end{pmatrix}.$$

The integration constants c_1 and c_2 are defined by the system $\mathbf{V}\mathbf{c} = \mathbf{x}_0$, which for this example is

$$\begin{pmatrix} -1 & 2 \\ 1 & 3 \end{pmatrix} \begin{pmatrix} c_1 \\ c_2 \end{pmatrix} = \begin{pmatrix} 0 \\ -4 \end{pmatrix}.$$

Solving the above equations reveals $c_1 = -8/5$ and $c_2 = -4/5$. The final solution to this systems of ODEs is

$$\mathbf{x}(t) = -\frac{8}{5} \begin{pmatrix} -1 \\ 1 \end{pmatrix} e^{-t} - \frac{4}{5} \begin{pmatrix} 2 \\ 3 \end{pmatrix} e^{4t}.$$

This solution requires the matrix \mathbf{A} be perfect and therefore have a complete set of eigenvectors.

The function $f(t) = e^{\lambda t}$ is an *eigenfunction* of the derivative operator, i.e.

$$\frac{d}{dt} f(t) = \lambda e^{\lambda t} = \lambda f(t)$$

. The solution of a system of linear ODEs is the product of the eigenvectors of \mathbf{A} and the eigenfunctions of $\frac{d\mathbf{x}}{dt}$.

18.3.2 Stability of Linear ODEs

The eigenvalues of \mathbf{A} are sufficient to tell if the system $\frac{d\mathbf{x}}{dt} = \mathbf{A}\mathbf{x}$ is stable. For a system of linear ODEs to be stable, all eigenvalues of \mathbf{A} must be nonpositive. If the eigenvalues are all negative, each term $e^{\lambda_i t}$ goes to zero at long times, so all variables in the system go to zero. If any of the eigenvalues are zero, the system is still stable (provided all other eigenvalues are negative), but the system will go to a constant value $c_i \mathbf{v}_i$, where \mathbf{v}_i is the eigenvector associated with the zero eigenvalue.

18.3.3 Positive Definite Matrices

A symmetric matrix \mathbf{A} is *positive definite* (p.d.) if $\mathbf{x}^T \mathbf{A} \mathbf{x} > 0$ for all nonzero vectors \mathbf{x} . If a matrix \mathbf{A} satisfies the slightly relaxed requirement that $\mathbf{x}^T \mathbf{A} \mathbf{x} \geq 0$ for all nonzero \mathbf{x} , we say that \mathbf{A} is *positive semi-definite* (p.s.d.).

Knowing that a matrix is positive (semi-)definite is useful for quadratic programming problems like the Support Vector Machine classifier. The quadratic function $f(\mathbf{x}) = \mathbf{x}^T \mathbf{Q} \mathbf{x}$ is convex if and only if the matrix \mathbf{Q} is positive semi-definite. For optimization problems like quadratic programs, the convexity of the objective function has enormous implications. Convex quadratic programs must only have global optima, making them easy to solve using numerical algorithms.

Determining if a matrix is positive (semi-)definite can be difficult unless we use eigenvalues. Any matrix with only positive eigenvalues is positive definite, and any matrix with only nonnegative eigenvalues is positive semi-definite. For example, consider the 2×2 identity matrix

$$\mathbf{I} = \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix}.$$

The product $\mathbf{x}^T \mathbf{I} \mathbf{x}$ is

$$\begin{pmatrix} x_1 & x_2 \end{pmatrix} \begin{pmatrix} 1 & 0 \\ 0 & 1 \end{pmatrix} \begin{pmatrix} x_1 \\ x_2 \end{pmatrix} = x_1^2 + x_2^2.$$

Since $x_1^2 + x_2^2$ is greater than zero for all nonzero inputs x_1 and x_2 , the matrix \mathbf{I} is positive definite and all its eigenvalues should be positive. Indeed, the eigenvalues for the identity matrix are $\lambda_1 = \lambda_2 = 1$.

As another example, consider the matrix

$$\mathbf{A} = \begin{pmatrix} 1 & -2 \\ -2 & 1 \end{pmatrix}.$$

The product $\mathbf{x}^T \mathbf{A} \mathbf{x} = x_1^2 - 4x_1x_2 + x_2^2$, which is not always positive. When $x_1 = x_2 = 1$, we see that $x_1^2 - 4x_1x_2 + x_2^2 = -2$. We know that \mathbf{A} is not positive definite (or even

Remember that a matrix \mathbf{A} is symmetric if $\mathbf{A} = \mathbf{A}^T$.

If \mathbf{Q} is positive definite (rather than just positive semi-definite) then $\mathbf{x}^T \mathbf{Q} \mathbf{x}$ is strictly convex.

positive semi-definite), so \mathbf{A} should have at least one negative eigenvalue. As expected, the eigenvalues for \mathbf{A} are $\lambda_1 = 3$ and $\lambda_2 = -1$.

18.3.4 Network Centrality

Networks are represented by collections of *nodes* connected by *edges*. When analyzing a network, it is common to ask which node occupies the most important position in the network. For example, which airport would cause the most problems if closed due to weather? In biological networks, the importance or *centrality* of an enzyme is used to prioritize drug targets.

We can quantify the centrality of each node in a network using random walks. We start by choosing a random node in the network. Then we randomly choose one of the edges connected to the node and travel to a new node. This process repeats again and again as we randomly traverse nodes and edges. The centrality of each node is proportional to the number of times we visit the node during the random walk. In the airport analogy, randomly traveling to cities across the country means you will frequently visit major hub airports.

Measuring centrality by random walks is easy to understand but impractical for large networks. It could take millions of steps to repeatedly reach all the nodes in networks with only a few thousand of nodes. In practice, we use eigenvectors to calculate centrality for networks. We begin by constructing an *adjacency matrix* for the network. The adjacency matrix encodes the connections (edges) between the nodes. The adjacency matrix is square with rows and columns corresponding to nodes in the network. The (i,j) entry in the network is set to 1 if there is an edge connecting node i with node j . Otherwise, the entries are zero. Note that we only consider direct connections. If node 1 is connected to node 2 and node 2 is connected to node 3, we do not connect nodes 1 and 3 unless there is a direct edge between them. Also, no node is connected to itself, so the diagonal elements are always zero.

Consider the four node network shown in Figure 18.1. The four node network has the following adjacency matrix.

$$\begin{array}{c} \begin{array}{c} A \\ B \\ C \\ D \end{array} \begin{pmatrix} & A & B & C & D \\ 0 & 1 & 1 & 0 \\ 1 & 0 & 1 & 1 \\ 1 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 \end{pmatrix} \end{array}$$

To identify the most central node in the network, we find the eigenvector that corresponds to the largest eigenvalue (λ_{\max}). For the above adjacency matrix,




Figure 18.1: Sample network with four nodes and four edges.

$\lambda_{\max} \approx 2.2$, and the associated eigenvector is

$$\mathbf{v}_{\max} = \begin{pmatrix} 0.52 \\ 0.61 \\ 0.52 \\ 0.28 \end{pmatrix}.$$

The entries in the eigenvector \mathbf{v}_{\max} are called the *eigencentality* scores. The largest entry corresponds to the most central node, and the smallest entry is associated with the least central node. We see that node *B* (entry 2) is most central in Figure 18.1 and node *D* (entry 4) is least central.

In simple networks like Figure 18.1, the most central node also has the most direct connections. This is not always the case. Eigencentality considers not only the number of connections but also their importance. Each edge is weighted by the centrality of the nodes it connect. Connections from more central nodes are more important, just as flights between major hub cities usually have the highest passenger volumes. Eigencentality has numerous applications including web searching. Google uses a modified version of centrality (called PageRank) to determine which results should be displayed first to users.

Centrality requires only the leading eigenvector of a network's adjacency matrix. Power Iteration (section 18.2) can find the leading eigenvector efficiently for large networks.

18.4 Geometric Interpretation of Eigenvalues

Consider a matrix $\mathbf{A} \in \mathbb{R}^2$ with eigenvalues λ_1 and λ_2 and corresponding eigenvectors \mathbf{v}_1 and \mathbf{v}_2 . Let's take a vector \mathbf{x} and decompose it over the eigenvectors.

$$\mathbf{x} = a_1\mathbf{v}_1 + a_2\mathbf{v}_2$$

We can represent the vector \mathbf{x} by plotting it; however, instead of using the normal Cartesian unit vectors as axes, we will use the eigenvectors. The values of \mathbf{x} along the "eigenaxes" are a_1 and a_2 . What happens when we multiply \mathbf{x} and \mathbf{A} ?

$$\mathbf{Ax} = \mathbf{A}(a_1\mathbf{v}_1 + a_2\mathbf{v}_2) = \lambda_1 a_1 \mathbf{v}_1 + \lambda_2 a_2 \mathbf{v}_2$$

Visually, multiplying by \mathbf{A} scales the values along the eigenvector axes. The scaling factor along each axis is the corresponding eigenvalue. To quantify the overall effect of multiplying by the matrix \mathbf{A} , we can compare the areas swept out by the vector \mathbf{x} before and after multiplication. The area is simply the product of the values along each axis, so the ratio becomes

$$\frac{\text{area}(\mathbf{Ax})}{\text{area}(\mathbf{x})} = \frac{\lambda_1 a_1 \lambda_2 a_2}{a_1 a_2} = \lambda_1 \lambda_2.$$

We can repeat the same calculation in three dimensions by looking at the ratio of the volume before and after multiplying by the matrix \mathbf{A} .

$$\frac{\text{volume}(\mathbf{Ax})}{\text{volume}(\mathbf{x})} = \frac{\lambda_1 a_1 \lambda_2 a_2 \lambda_3 a_3}{a_1 a_2 a_3} = \lambda_1 \lambda_2 \lambda_3$$

In general, the product of the eigenvalues of a matrix describe the overall effect of multiplying a vector by the matrix. The product of the eigenvalues of a matrix \mathbf{A} is called the *determinant* of \mathbf{A} , or $\det(\mathbf{A})$.

$$\det(\mathbf{A}) = \lambda_1 \lambda_2 \cdots \lambda_n$$

If the determinant of a matrix is large, multiplying a vector by the matrix enlarges the volume swept out by the vector. If the determinant is small, the volume contracts.

18.5 Properties of the Determinant

The determinant is a powerful property of a matrix. Determinants can be easily calculated for a matrix and contain useful information about the matrix.

Let's say a vector $\mathbf{y} = \mathbf{Ax}$. We know the determinant of the matrix \mathbf{A} is the ratio of the volumes of \mathbf{x} and \mathbf{y} .

$$\det(\mathbf{A}) = \frac{\text{volume}(\mathbf{Ax})}{\text{volume}(\mathbf{x})} = \frac{\text{volume}(\mathbf{y})}{\text{volume}(\mathbf{x})}$$

If the inverse of \mathbf{A} exists, we know that $\mathbf{x} = \mathbf{A}^{-1}\mathbf{y}$. Thus, the determinant of the inverse matrix \mathbf{A}^{-1} is

$$\det(\mathbf{A}^{-1}) = \frac{\text{volume}(\mathbf{A}^{-1}\mathbf{y})}{\text{volume}(\mathbf{y})} = \frac{\text{volume}(\mathbf{x})}{\text{volume}(\mathbf{y})} = \frac{1}{\det(\mathbf{A})}.$$

The determinant of \mathbf{A}^{-1} is the inverse of the determinant of \mathbf{A} . If the determinant of a matrix is zero, this property indicates there is a problem with the inverse of the matrix.

$$\det(\mathbf{A}) = 0 \Rightarrow \det(\mathbf{A}^{-1}) = \frac{1}{\det(\mathbf{A})} = \frac{1}{0} \rightarrow \text{undefined}$$

Although we won't prove it in this course, **a matrix has an inverse if and only if the determinant of the matrix is nonzero**. The following statements are, in fact, equivalent for a square matrix \mathbf{A} :

Remember that the volume we're discussing here is the volume when a vector is plotted using the matrix's eigenvectors as axes.

In MATLAB, the function `det` calculates the determinant of a matrix.

- \mathbf{A} can be transformed into the identity matrix by elementary row operations.
- The system $\mathbf{Ax} = \mathbf{y}$ is solvable and has a unique solution.
- \mathbf{A} is full rank.
- \mathbf{A}^{-1} exists.
- $\det(\mathbf{A}) \neq 0$.

A matrix with a determinant equal to zero has a geometric interpretation. Remember that the determinant is the product of the eigenvalues. It is the product of the scaling factors of the matrix along each eigenvector. If one of the eigenvalues is zero, we are missing information about how the matrix scales along at least one eigenvector. Our knowledge of the transformation is incomplete, which is why we cannot find a unique solution for the corresponding linear system.

Using the determinant we can concisely state our last field axiom. Recall that for scalars we required a multiplicative inverse exist for any nonzero member of the field, i.e.

For all scalars $a \neq 0$ there exists a^{-1} such that $aa^{-1} = 1$.

For vector spaces, we have the following:

For all square matrices \mathbf{A} where $\det(\mathbf{A}) \neq 0$,
there exists \mathbf{A}^{-1} such that $\mathbf{AA}^{-1} = \mathbf{I} = \mathbf{A}^{-1}\mathbf{A}$.

Chapter 19

Matrix Decompositions

19.1 Eigendecomposition

Let's discuss a square, $n \times n$ matrix \mathbf{A} . Provided \mathbf{A} is not defective, it has n linearly independent eigenvectors which we will call $\mathbf{v}_1, \dots, \mathbf{v}_n$. The eigenvectors are linearly independent and therefore form a basis for \mathbb{R}^n (an *eigenbasis*). We said in the last chapter that any vector \mathbf{x} can be decomposed onto the eigenbasis by finding coefficients a_1, \dots, a_n such that

$$\mathbf{x} = a_1 \mathbf{v}_1 + a_2 \mathbf{v}_2 + \dots + a_n \mathbf{v}_n.$$

Multiplying the vector \mathbf{x} by the matrix \mathbf{A} is equivalent to scaling each term in the decomposition by the corresponding eigenvalue (λ_i).

$$\mathbf{A}\mathbf{x} = a_1 \lambda_1 \mathbf{v}_1 + a_2 \lambda_2 \mathbf{v}_2 + \dots + a_n \lambda_n \mathbf{v}_n$$

We can think of matrix multiplication as a transformation with three steps.

1. Decompose the input vector onto the eigenbasis of the matrix.
2. Scale each term in the decomposition by the appropriate eigenvalue.
3. Reassemble, or “un-decompose” the output vector.

Each of these steps can be represented by a matrix operation. First, we collect the n eigenvalue into a matrix \mathbf{V} .

$$\mathbf{V} = (\mathbf{v}_1 \ \mathbf{v}_2 \ \dots \ \mathbf{v}_n)$$

Each column in the matrix \mathbf{V} is an eigenvector of the matrix \mathbf{A} . To decompose the vector \mathbf{x} onto the columns of \mathbf{V} we find the coefficients a_1, \dots, a_n by solving the linear system

$$\mathbf{V}\mathbf{a} = \mathbf{x}$$

where \mathbf{a} is a vector holding the coefficients a_1, \dots, a_n . The matrix \mathbf{V} is square and has linearly independent columns (the eigenvectors of \mathbf{A}), so its inverse exists. The coefficients for decomposing the vector \mathbf{x} onto the eigenbasis of the matrix \mathbf{A} are

$$\mathbf{a} = \mathbf{V}^{-1}\mathbf{x}.$$

If the inverse matrix \mathbf{V}^{-1} decomposes a vector into a set of coefficients \mathbf{a} , then multiplying the coefficients vector \mathbf{a} by the original matrix must reassemble the vector \mathbf{x} . Looking back at the three steps we defined above, we can use multiplication by \mathbf{V}^{-1} to complete step 1 and multiply by \mathbf{V} to perform step 3. For step 2, we need to scale the individual coefficients by the appropriate eigenvalues. We define a scaling matrix $\mathbf{\Lambda}$ as a diagonal matrix of the eigenvalues:

$$\mathbf{\Lambda} = \begin{pmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_n \end{pmatrix}.$$

Notice that

$$\mathbf{\Lambda}\mathbf{a} = \begin{pmatrix} \lambda_1 a_1 \\ \lambda_2 a_2 \\ \vdots \\ \lambda_n a_n \end{pmatrix}$$

so the matrix $\mathbf{\Lambda}$ scales the i th entry of the input vector by the i th eigenvalue.

We now have matrix operations for decomposing onto an eigenbasis (\mathbf{V}^{-1}), scaling by eigenvalues ($\mathbf{\Lambda}$), and reassembling the output vector (\mathbf{V}). Putting everything together, we see that matrix multiplication can be expressed as an *eigendecomposition* by

$$\mathbf{A}\mathbf{x} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}\mathbf{x}.$$

Equivalently, we can say that the matrix \mathbf{A} itself can be as the product of three matrices ($\mathbf{A} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}$) if \mathbf{A} has a complete set of eigenvectors. There are two ways to interpret the dependence on a complete set of eigenvectors. Viewed technically, the matrix \mathbf{V} can only be inverted if it is full rank, so \mathbf{V}^{-1} does not exist if one or more eigenvectors is missing. More intuitively, the eigendecomposition defines

In other words, if \mathbf{V}^{-1} decomposes a vector, $(\mathbf{V}^{-1})^{-1} = \mathbf{V}$ must undo the decomposition.

We use the uppercase Greek lambda ($\mathbf{\Lambda}$) to denote the matrix of eigenvalues λ_i (lowercase lambda).

Eigendecomposition is the last time we will use the prefix "eigen-". Feel free to use it on other everyday words to appear smarter.

a unique mapping between the input and output vectors. Uniqueness requires a basis, since a vector decomposition is only unique if the set of vectors form a basis. If the matrix \mathbf{A} is defective, its eigenvectors do not form an eigenbasis and there cannot be a unique mapping between inputs and outputs.

19.2 Singular Value Decomposition

The eigendecomposition is limited to square matrices with a complete set of eigenvectors. However, the idea that matrices can be factored into three operations (decomposition, scaling, and reassembly) generalizes to all matrices, even non-square matrices. The generalized equivalent of the eigendecomposition is called the *Singular Value Decomposition*, or (SVD).

Singular Value Decomposition. Any $m \times n$ matrix \mathbf{A} can be factored into the product of three matrices

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

where

- \mathbf{U} is an orthogonal $m \times m$ matrix.
- $\mathbf{\Sigma}$ is a diagonal $m \times n$ matrix with nonzero entries.
- \mathbf{V} is an orthogonal $n \times n$ matrix.

The square matrices \mathbf{U} and \mathbf{V} are *orthogonal*, i.e. their columns form an orthonormal set of basis vectors. As we discussed previously, the inverse of an orthogonal matrix is simply its transpose, so $\mathbf{U}^{-1} = \mathbf{U}^T$ and $\mathbf{V}^{-1} = \mathbf{V}^T$. The \mathbf{V}^T term in the decomposition has the same role as the \mathbf{V}^{-1} matrix in an eigendecomposition — projection of the input vector onto a new basis. The matrix \mathbf{U} in SVD reassembles the output vector analogous to the vector \mathbf{V} in an eigendecomposition.

The matrix $\mathbf{\Sigma}$ is diagonal but not necessarily square. It has the same dimensions as the original matrix \mathbf{A} . For a 3×5 matrix, the $\mathbf{\Sigma}$ has the form

$$\mathbf{\Sigma} = \begin{pmatrix} \sigma_1 & 0 & 0 & 0 & 0 \\ 0 & \sigma_2 & 0 & 0 & 0 \\ 0 & 0 & \sigma_3 & 0 & 0 \end{pmatrix}.$$

If the matrix \mathbf{A} was 5×3 , we would have

$$\mathbf{\Sigma} = \begin{pmatrix} \sigma_1 & 0 & 0 \\ 0 & \sigma_2 & 0 \\ 0 & 0 & \sigma_3 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}$$

If the entries in \mathbf{A} were complex numbers, the matrices \mathbf{U} and \mathbf{V} would be *unitary*. The inverse of a unitary matrix is the complex conjugate of the matrix transpose.

The elements along the diagonal of Σ are called *singular values*. If \mathbf{A} is an $m \times n$ matrix, the maximum number of nonzero singular values is $\min\{m, n\}$. They are the analogues of eigenvalues for non-square matrices. However, the singular values for a square matrix are not equal to the eigenvalues of the same matrix. Singular values are always nonnegative. If we arrange Σ such that the singular values are in descending order, the SVD of a matrix is unique.

The columns in \mathbf{U} and \mathbf{V} are called the left and right *singular vectors*, respectively. Just as there is a relationship between eigenvalues and eigenvectors, the columns in \mathbf{U} and \mathbf{V} are connected by the singular values in Σ . If $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T$, then

$$\mathbf{A}\mathbf{v}_i = \sigma_i\mathbf{u}_i$$

and

$$\mathbf{A}^T\mathbf{u}_i = \sigma_i\mathbf{v}_i$$

where \mathbf{v}_i is i th right singular vector (the i th column in \mathbf{V}); \mathbf{u}_i is the i th left singular vector (the i th column in \mathbf{U}); and σ_i is the i th singular value (the i th nonzero on the diagonal of Σ).

19.3 Applications of the SVD

19.3.1 Rank of a matrix

The rank of a matrix \mathbf{A} is equal to the number of nonzero singular values (the number of nonzero values along the diagonal of Σ). This is true for both square and nonsquare matrices. Notice that the way we defined the diagonal matrix Σ implies that the number of singular values must be at most $\min\{m, n\}$ for an $m \times n$ matrix. This requirement agrees with our knowledge that $\text{rank}(\mathbf{A}) \leq \min\{m, n\}$.

19.3.2 The matrix pseudoinverse

Our definition of a matrix inverse applies only to square matrices. For nonsquare matrices we can use the SVD to construct a pseudoinverse. We represent the pseudoinverse of a matrix \mathbf{A} as \mathbf{A}^+ . We simply reverse and invert the factorization of \mathbf{A} , i.e.

$$\mathbf{A}^+ = (\mathbf{V}^T)^{-1}\Sigma^+\mathbf{U}^{-1}$$

We can simplify this expression with knowledge that \mathbf{V} and \mathbf{U} are orthogonal, so $(\mathbf{V}^T)^{-1} = (\mathbf{V}^T)^T = \mathbf{V}$ and $\mathbf{U}^{-1} = \mathbf{U}^T$. Thus

$$\mathbf{A}^+ = \mathbf{V}\Sigma^+\mathbf{U}^T.$$

The matrix Σ^+ is the pseudoinverse of the diagonal matrix Σ . This is simply the transpose of Σ where each entry on the diagonal is replaced by its multiplicative inverse. For example, a 3×5 matrix Σ :

$$\Sigma = \begin{pmatrix} \sigma_1 & 0 & 0 & 0 & 0 \\ 0 & \sigma_2 & 0 & 0 & 0 \\ 0 & 0 & \sigma_3 & 0 & 0 \end{pmatrix}$$

the pseudoinverse Σ^+ is

$$\Sigma^+ = \begin{pmatrix} 1/\sigma_1 & 0 & 0 \\ 0 & 1/\sigma_2 & 0 \\ 0 & 0 & 1/\sigma_3 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}.$$