

Matrix Decompositions

BIOE 210

Matrix multiplication using the eigenbasis

Let's consider a square $n \times n$ matrix \mathbf{A} with a complete set of eigenvectors $\mathbf{v}_1, \mathbf{v}_2, \dots, \mathbf{v}_n$.

The eigenvectors form a basis (the *eigenbasis*), so we can decompose another vector \mathbf{x} onto them.

$$\mathbf{x} = a_1 \mathbf{v}_1 + \dots + a_n \mathbf{v}_n$$

We can use the decomposition of \mathbf{x} to perform matrix multiplication by \mathbf{A} .

$$\mathbf{Ax} = a_1 \lambda_1 \mathbf{v}_1 + \dots + a_n \lambda_n \mathbf{v}_n$$

Matrix multiplication as a multistep transformation

Using the eigenbasis we can break down matrix multiplication into three steps:

1. Decompose \mathbf{x} onto the eigenbasis.
2. Scale each term by the respective eigenvalue.
3. Reassemble the decomposed vectors into the output vector.

Step 1: Decompose \mathbf{x} onto the eigenbasis

We already know how to decompose a vector onto a basis. First we assemble a matrix using the basis vectors as columns.

$$\mathbf{V} = (\mathbf{v}_1 \ \mathbf{v}_2 \ \cdots \ \mathbf{v}_n)$$

Then we find the coefficients \mathbf{a} by solving the linear system

$$\mathbf{V}\mathbf{a} = \mathbf{x}$$

This has a unique solution

$$\mathbf{a} = \mathbf{V}^{-1}\mathbf{x}$$

Step 2: Scale the coefficients by the eigenvalues

Multiplying by the matrix **A** scales each coefficient in our decomposition by the corresponding eigenvalue:

$$a_1 \rightarrow \lambda_1 a_1$$

$$a_2 \rightarrow \lambda_2 a_2$$

$$\vdots$$

$$a_n \rightarrow \lambda_n a_n$$

We can scale all of the coefficients at once using a matrix with the eigenvalues along the diagonal:

$$\mathbf{\Lambda} = \begin{pmatrix} \lambda_1 & 0 & \cdots & 0 \\ 0 & \lambda_2 & \cdots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \cdots & \lambda_n \end{pmatrix}, \quad \mathbf{\Lambda a} = \begin{pmatrix} \lambda_1 a_1 \\ \lambda_2 a_2 \\ \vdots \\ \lambda_n a_n \end{pmatrix}$$

Step 3: Reassembling the output vector

We can also use matrix multiplication to reassemble a decomposed vector into the output vector.

Notice that the inverse matrix \mathbf{V}^{-1} decomposes a vector.

$$\mathbf{a} = \mathbf{V}^{-1} \mathbf{x}$$

The original matrix undoes this operation since it is the inverse of the inverse.

$$\begin{aligned}\mathbf{Va} &= \mathbf{VV}^{-1} \mathbf{x} \\ &= \mathbf{x}\end{aligned}$$

Reassembling the decomposed vector is the same as multiplying by the matrix \mathbf{V} .

Putting it all together

We have broken down matrix multiplication into three sequential matrix operations:

$$\begin{aligned}\mathbf{Ax} &= \underbrace{(\text{reassembly})}_{\mathbf{V}} \underbrace{(\text{scaling})}_{\mathbf{\Lambda}} \underbrace{(\text{decomposition})}_{\mathbf{V}^{-1}} \mathbf{x} \\ &= \mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}\mathbf{x}\end{aligned}$$

The matrix \mathbf{A} can therefore be decomposed into the product of three matrices, $\mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}$.

This matrix decomposition is called the *eigendecomposition* of \mathbf{A} .

Limitations of the eigendecomposition

The eigendecomposition $\mathbf{A} = \mathbf{V}\mathbf{\Lambda}\mathbf{V}^{-1}$ only works when the matrix \mathbf{A} meets certain requirements:

- ▶ \mathbf{A} must be square.
- ▶ \mathbf{A} must have a complete set of eigenvectors.
- ▶ \mathbf{V}^{-1} must exist.¹

There is a more general form of matrix decomposition that relaxes these requirements.

¹The inverse \mathbf{V}^{-1} will always exist when \mathbf{A} has a complete set of eigenvectors since the eigenvectors are linearly independent and therefore \mathbf{V} is full rank.

The Singular Value Decomposition (SVD)

Any $m \times n$ matrix \mathbf{A} can be decomposed into the product of three matrices

$$\mathbf{A} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T$$

where

- ▶ \mathbf{U} is an orthogonal $m \times m$ matrix.
- ▶ $\mathbf{\Sigma}$ is a diagonal $m \times n$ matrix with nonzero entries.
- ▶ \mathbf{V} is an orthogonal $n \times n$ matrix.

Parts of the SVD

The matrices **U** and **V** are *orthogonal* matrices, meaning their columns are an orthonormal set of basis vectors. This also means that

$$\mathbf{U}^{-1} = \mathbf{U}^T, \quad \mathbf{V}^{-1} = \mathbf{V}^T$$

The columns in **U** and **V** are called the left and right *singular vectors*. The singular vectors have similar properties to eigenvectors:

$$\mathbf{A}\mathbf{v}_i = \sigma_i\mathbf{u}_i$$

and

$$\mathbf{A}^T\mathbf{u}_i = \sigma_i\mathbf{v}_i$$

where the scalars σ_i are the *singular values*.

Parts of the SVD

The matrix Σ is a diagonal $m \times n$ matrix. The entries along the diagonal are the singular values.

Singular values are always non-negative. They are similar to eigenvalues for square matrices, but they are not the same.²

If we arrange Σ such that the singular values are in descending order, the SVD of a matrix is unique.

²The singular values of a square matrix are the squares of the eigenvalues of the matrix.

Applications of the SVD

We will cover several examples of the SVD, including:

- ▶ Calculating the rank of a matrix. (textbook)
- ▶ Finding a “pseudoinverse” of non-square matrices. (textbook)
- ▶ Approximating matrices with simpler matrices.
- ▶ Image compression.