


Neural Networks: Perceptrons

Spring 2021


The artificial neuron

A neuron connects a series on inputs (dendrites) to an output (axon).



The artificial neuron


A neuron connects a series on inputs (dendrites) to an output (axon).



Our model needs to include two processes:

1. The cell body (soma) combines all of the n inputs.
2. If the combined input exceeds a threshold, the output fires.


Modeling the artificial neuron



Let's model the combined input z as a linear combination of the inputs.

$$z = w_1x_1 + w_2x_2 + \cdots + w_nx_n = \mathbf{w} \cdot \mathbf{x}$$

Modeling the artificial neuron



Let's model the combined input z as a linear combination of the inputs.

$$z = w_1x_1 + w_2x_2 + \cdots + w_nx_n = \mathbf{w} \cdot \mathbf{x}$$

For now, let's assume the neuron "fires" based on the sign of z :

$$y = \begin{cases} +1, & \mathbf{w} \cdot \mathbf{x} > 0 \\ -1, & \mathbf{w} \cdot \mathbf{x} < 0 \end{cases}$$

Learning to classify

Let's assume we have n training points (\mathbf{x}_i, y_i) . How can we use these data to find the weights \mathbf{w} so the neuron fires only when $y_i = +1$?

Learning to classify

Let's assume we have n training points (\mathbf{x}_i, y_i) . How can we use these data to find the weights \mathbf{w} so the neuron fires only when $y_i = +1$?


The Perceptron Update Algorithm

1. Choose a random vector of initial weights $\mathbf{w}^{(0)}$ and a step size α .
2. For each training point (\mathbf{x}_i, y_i) , update the weights by


$$\mathbf{w}^{(k+1)} = \begin{cases} \mathbf{w}^{(k)} + \alpha y_i \mathbf{x}_i, & \text{if } \mathbf{x}_i \text{ is misclassified} \\ \mathbf{w}^{(k)}, & \text{otherwise} \end{cases}$$

3. Repeat until all points are classified correctly.

Does it work?




Heart attack classification by SVM.




Heart attack classification by a perceptron.

Why does it work?



Point \mathbf{x} is $+1$ and is
therefore misclassified by
 $\mathbf{w}^{(k)}$.

Why does it work?




Point \mathbf{x} is $+1$ and is therefore misclassified by $\mathbf{w}^{(k)}$.

The weight is updated by

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \alpha \mathbf{x},$$

tilting \mathbf{w} toward \mathbf{x} .

Why does it work?




Point x is $+1$ and is therefore misclassified by $\mathbf{w}^{(k)}$.

The weight is updated by

$$\mathbf{w}^{(k+1)} = \mathbf{w}^{(k)} + \alpha \mathbf{x},$$

tilting \mathbf{w} toward x .



The new hyperplane is rotated toward x so the point will eventually be on the correct side.

Wait, what happened to the intercept?

The SVM problem used a linear classifier

$$y = \begin{cases} +1, & \mathbf{a} \cdot \mathbf{x} \geq b \\ -1, & \mathbf{a} \cdot \mathbf{x} \leq b \end{cases}$$

but our perceptron fires using the rule

$$y = \begin{cases} +1, & \mathbf{w} \cdot \mathbf{x} > 0 \\ -1, & \mathbf{w} \cdot \mathbf{x} < 0 \end{cases}$$

Does this mean the perceptron hyperplane always passes through the origin?

Wait, what happened to the intercept?

The SVM problem used a linear classifier

$$y = \begin{cases} +1, & \mathbf{a} \cdot \mathbf{x} \geq b \\ -1, & \mathbf{a} \cdot \mathbf{x} \leq b \end{cases}$$

but our perceptron fires using the rule

$$y = \begin{cases} +1, & \mathbf{w} \cdot \mathbf{x} > 0 \\ -1, & \mathbf{w} \cdot \mathbf{x} < 0 \end{cases}$$

Does this mean the perceptron hyperplane always passes through the origin?

No. We use a common ML trick to move the *bias* (intercept) into the weight vector and expand \mathbf{x} with a dummy dimension containing 1.


$$\mathbf{w} \cdot \mathbf{x} = b \Leftrightarrow \begin{pmatrix} w_1 \\ w_2 \\ -b \end{pmatrix} \cdot \begin{pmatrix} x_1 \\ x_2 \\ 1 \end{pmatrix} = 0$$

Example: Training a 2D Perceptron

x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1
-0.5	-0.7	-1
0.4	-0.3	-1


$k = 0$, accuracy = 4/6

$$\mathbf{w}^{(0)} = (-0.2, 0.05, 0.03)^\top$$



Example: Training a 2D Perceptron

x_1	x_2	y	
1	0.9	+1	←
-1	0.3	+1	
0	0.5	+1	
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	



$$k = 0, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(0)} = (-0.2, 0.05, 0.03)^\top$$

$$\mathbf{w}^{(0)} \cdot \mathbf{x} = -0.125 \quad (\text{mismatch})$$

$$\begin{aligned}\mathbf{w}^{(1)} &= \mathbf{w}^{(0)} + 0.1(+1)\mathbf{x} \\ &= (-0.1, 0.14, 0.13)^\top\end{aligned}$$

Example: Training a 2D Perceptron

x_1	x_2	y	
1	0.9	+1	←
-1	0.3	+1	
0	0.5	+1	
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	


$$k = 0, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(0)} = (-0.2, 0.05, 0.03)^T$$

$$\mathbf{w}^{(0)} \cdot \mathbf{x} = -0.125 \quad (\text{mismatch})$$

$$\mathbf{w}^{(1)} = \mathbf{w}^{(0)} + 0.1(+1)\mathbf{x}$$

$$= (-0.1, 0.14, 0.13)^T$$



Example: Training a 2D Perceptron


x_1	x_2	y	
1	0.9	+1	
-1	0.3	+1	←
0	0.5	+1	
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	

$$k = 1, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(1)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(1)} \cdot \mathbf{x} = 0.272 \quad (\text{match})$$

$$\mathbf{w}^{(2)} = \mathbf{w}^{(1)}$$



Example: Training a 2D Perceptron


x_1	x_2	y	
1	0.9	+1	
-1	0.3	+1	←
0	0.5	+1	
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	

$k = 1$, accuracy = 4/6

$$\mathbf{w}^{(1)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(1)} \cdot \mathbf{x} = 0.272 \quad (\text{match})$$

$$\mathbf{w}^{(2)} = \mathbf{w}^{(1)}$$



Example: Training a 2D Perceptron


x_1	x_2	y	
1	0.9	+1	
-1	0.3	+1	
0	0.5	+1	←
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	

$$k = 2, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(2)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(2)} \cdot \mathbf{x} = 0.2 \quad (\text{match})$$

$$\mathbf{w}^{(3)} = \mathbf{w}^{(2)}$$



Example: Training a 2D Perceptron


x_1	x_2	y	
1	0.9	+1	
-1	0.3	+1	
0	0.5	+1	←
1	-1	-1	
-0.5	-0.7	-1	
0.4	-0.3	-1	

$$k = 2, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(2)} = (-0.1, 0.14, 0.13)^T$$


$$\mathbf{w}^{(2)} \cdot \mathbf{x} = 0.2 \quad (\text{match})$$

$$\mathbf{w}^{(3)} = \mathbf{w}^{(2)}$$



Example: Training a 2D Perceptron

x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1 ←
-0.5	-0.7	-1
0.4	-0.3	-1



$$k = 3, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(3)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(3)} \cdot \mathbf{x} = -0.11 \quad (\text{match})$$

$$\mathbf{w}^{(4)} = \mathbf{w}^{(3)}$$

Example: Training a 2D Perceptron


x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1 ←
-0.5	-0.7	-1
0.4	-0.3	-1

$k = 3$, accuracy = 4/6

$$\mathbf{w}^{(3)} = (-0.1, 0.14, 0.13)^T$$


$$\mathbf{w}^{(3)} \cdot \mathbf{x} = -0.11 \quad (\text{match})$$

$$\mathbf{w}^{(4)} = \mathbf{w}^{(3)}$$



Example: Training a 2D Perceptron

x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1
-0.5	-0.7	-1
0.4	-0.3	-1



$$k = 4, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(4)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(4)} \cdot \mathbf{x} = 0.082 \quad (\text{mismatch})$$

$$\begin{aligned}\mathbf{w}^{(5)} &= \mathbf{w}^{(4)} + 0.1(-1)\mathbf{x} \\ &= (-0.05, 0.21, 0.03)^T\end{aligned}$$

Example: Training a 2D Perceptron

x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1
-0.5	-0.7	-1
0.4	-0.3	-1


$$k = 4, \quad \text{accuracy} = 4/6$$

$$\mathbf{w}^{(4)} = (-0.1, 0.14, 0.13)^T$$

$$\mathbf{w}^{(4)} \cdot \mathbf{x} = 0.082 \quad (\text{mismatch})$$


$$\mathbf{w}^{(5)} = \mathbf{w}^{(4)} + 0.1(-1)\mathbf{x}$$

$$= (-0.05, 0.21, 0.03)^T$$



Example: Training a 2D Perceptron

x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1
-0.5	-0.7	-1
0.4	-0.3	-1



$k = 5$, accuracy = 6/6

$$\mathbf{w}^{(5)} = (-0.05, 0.21, 0.03)^T$$

$$\mathbf{w}^{(5)} \cdot \mathbf{x} = -0.05 \quad (\text{match})$$

$$\mathbf{w}^{(6)} = \mathbf{w}^{(5)}$$

Example: Training a 2D Perceptron


x_1	x_2	y
1	0.9	+1
-1	0.3	+1
0	0.5	+1
1	-1	-1
-0.5	-0.7	-1
0.4	-0.3	-1

$k = 5$, accuracy = 6/6

$$\mathbf{w}^{(5)} = (-0.05, 0.21, 0.03)^T$$

$$\mathbf{w}^{(5)} \cdot \mathbf{x} = -0.05 \quad (\text{match})$$

$$\mathbf{w}^{(6)} = \mathbf{w}^{(5)}$$



Summary

- ▶ A perceptron is a simplistic model of a single neuron.
- ▶ A perceptron can learn to perform simple classification tasks using an update rule.

Summary

- ▶ A perceptron is a simplistic model of a single neuron.
- ▶ A perceptron can learn to perform simple classification tasks using an update rule.
- ▶ **Next time:** Imagine what a network of millions of perceptrons can learn!